

# The Rust Programming Language

Pramode C.E

March 2, 2017

*A language that doesn't affect the way you think about programming, is not worth knowing.*

Alan Perlis.

# Why is Rust so exciting?

- ▶ Safe low-level systems programming
- ▶ Memory safety without garbage collection
- ▶ High-level abstractions (influenced by statically typed functional programming languages like ML) without run-time overhead
- ▶ Concurrency without data races

## Why not just use C/C++?

*The second big thing each student should learn is: How can I avoid being burned by C's numerous and severe shortcomings? This is a development environment where only the paranoid can survive.*

<http://blog.regehr.org/archives/1393>

## A bit of Rust history

- ▶ Started by Graydon Hoare as a personal project in 2006
- ▶ Mozilla foundation started sponsoring Rust in 2010
- ▶ Rust 1.0 released in May, 2015
- ▶ Regular six week release cycles

# Core language features

- ▶ Memory safety without garbage collection
  - ▶ Ownership
  - ▶ Move Semantics
  - ▶ Borrowing and lifetimes
- ▶ Type Inference
- ▶ Algebraic Data Types (Sum and Product types)
- ▶ Exhaustive Pattern Matching
- ▶ Trait-based generics
- ▶ Iterators
- ▶ Zero Cost Abstractions
- ▶ Concurrency without data races
- ▶ Efficient C bindings, minimal runtime

# Structure of this workshop

- ▶ Understanding the problems with C/C++
- ▶ Understanding Ownership, Borrow, Move semantics and Lifetimes
- ▶ Other features (depending on availability of time)

# The Stack

```
// a0.c
```

```
int main()
{
    int a = 1, b = 2;
    return 0;
}
```



# The Stack

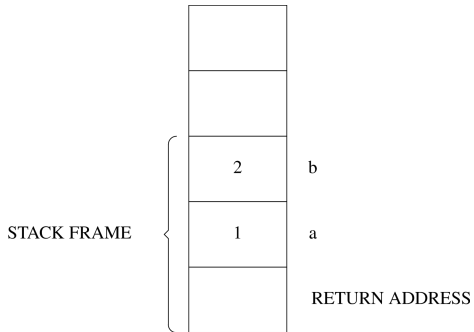


Figure 1: The C stack

# Buffer overflow

```
// a1.c
int main()
{
    int i=1;
    int a[4];
    int j=2;

    a[4] = 5; // bug
    a[5] = 6; // bug
    a[10000] = 7; // bug
}
```

# Buffer overflow

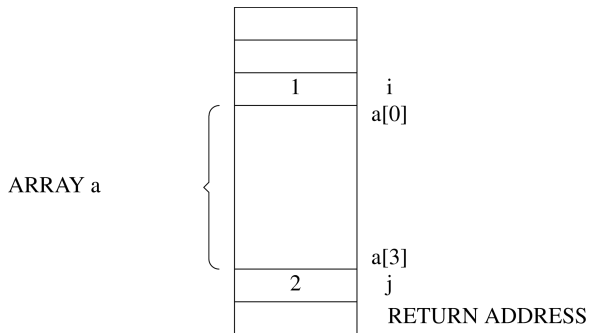


Figure 2: Buffer overflow diagram

# Pointers in C

```
// a2.c  
int main()  
{  
    int a = 10;  
    int *b;  
  
    b = &a;  
    *b = 20;  
}
```

# Pointers in C

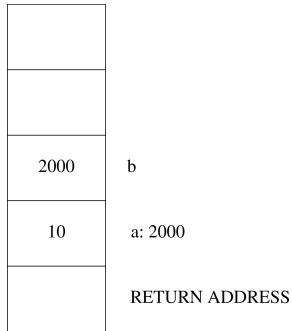


Figure 3: How pointer variables are stored in memory

## Stack Frames - deallocations and allocations

```
// a3.c  
void fun2() { int e=5, f=6; }  
void fun1() { int c=3, d=4; }  
  
int main()  
{  
    int a=1, b=2;  
    fun1(); fun2();  
    return 0;  
}
```

# Stack Frames - allocations and deallocations




Figure 4: Multiple stack frames

# Stack Frames - allocations and deallocations

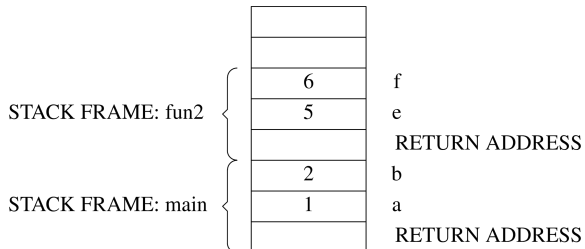


Figure 5: Multiple stack frames



# Dangling Pointers

```
// a4.c
void fun2() { int m = 1; int n = 2; }
int* fun1() {
    int *p; int q = 0;
    p = &q; return p; // bug
}

int main() {
    int *a, b; a = fun1();
    *a = 10; fun2();
    b = *a;
}
```

# Dangling pointers

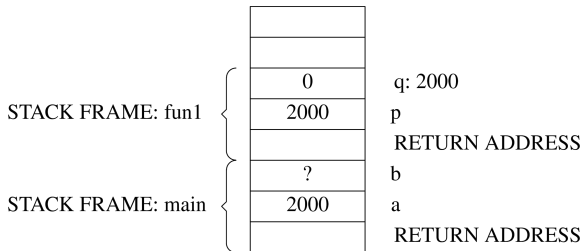


Figure 6: Dangling pointer

## Null pointer dereferencing

```
// a6.c
#include <strings.h>
#include <stdio.h>
int main()
{
    char s[100] = "hello";
    char *p;
    p = index(s, 'f');
    *p = 'a'; // bug!
    return 0;
}
```

## Heap allocation - malloc and free

```
// a7.c
#include <stdlib.h>
void fun()
{
    char *c;
    c = malloc(10*sizeof(char));
    /* do some stuff here */
    free(c);
}
int main()
{
    fun();
}
```

# Heap allocation - malloc and free

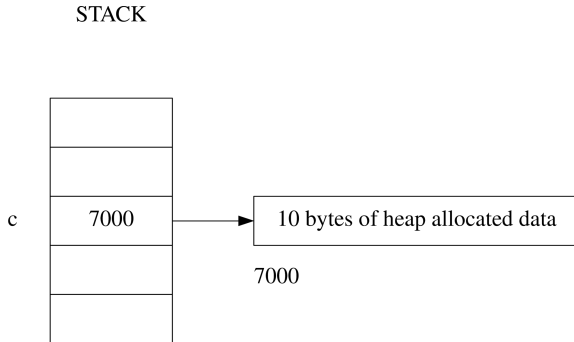


Figure 7: A stack location pointing to a heap location

# Memory leaks

```
// a8.c
#include <stdlib.h>
void fun()
{
    char *c;
    c = malloc(10*sizeof(char));
    /* do some stuff here */
}
int main()
{
    fun(); // bug! memory leak.
}
```

## Use-after-free

```
// a9.c
#include <stdlib.h>
void fun(char *t) {
    /* do some stuff here */
    free(t);
}
int main() {
    char *c;
    c = malloc(10 * sizeof(char));
    fun(c);
    c[0] = 'A'; //bug! user-after-free
}
```

## Double free

```
// a10.c
#include <stdlib.h>
void fun(char *t) {
    /* do some stuff here */
    free(t);
}
int main() {
    char *c;
    c = malloc(10 * sizeof(char));
    fun(c);
    free(c); //bug! double free
}
```



# Undefined behaviours and optimization

```
// a11.c
#include <limits.h>
#include <stdio.h>
// compile the code with optimization (-O3) and without
int main() {
    int c = INT_MAX;
    if (c+1 < c)
        printf("hello\n");
    printf("%d\n", c+1);
}
```

# Undefined behaviours

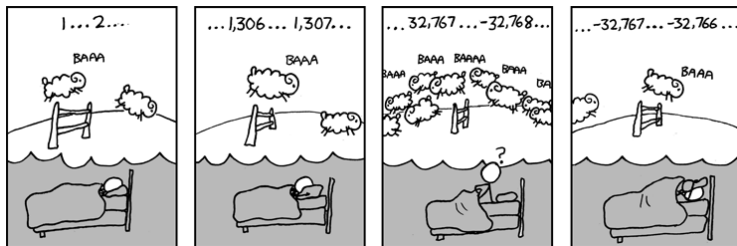


Figure 8: [https://www.explainxkcd.com/wiki/images/c/cc/cant\\_sleep.png](https://www.explainxkcd.com/wiki/images/c/cc/cant_sleep.png)

## Undefined behaviours

*When tools like the bounds checking GCC, Purify, Valgrind, etc. first showed up, it was interesting to run a random UNIX utility under them. The output of the checker showed that these utility programs, despite working perfectly well, executed a ton of memory safety errors such as use of uninitialized data, accesses beyond the ends of arrays, etc. Just running grep or whatever would cause tens or hundreds of these errors to happen.*

From: <http://blog.regehr.org/archives/226>

## Undefined behaviours

*More and more, I'm starting to wonder how safety-critical code can continue being written in C.*

A comment on: <http://blog.regehr.org/archives/232>

# Hello, world!

```
// a12.rs  
fn main() {  
    println!("hello, world!");  
}
```

Compile: `rustc a12.rs`

Run: `./a12`

# Type inference

```
// a12-1.rs
fn sqr(x: i32) -> i32 {
    let y = x * x; // type inferred
    y
}
fn main() {
    let t1 = sqr(10); // type inferred
    let t2:i32 = sqr(20);
    println!("sqr 10 = {}, sqr 20 ={}", t1, t2);
}
```

# Immutability

```
// a12-2.rs
fn main() {
    let x = 0; // x is immutable
    let mut y = 1;
    x = x + 1; // does not work
    y = y + 1;
}
```

# Scope

```
// a12-3.rs  
fn main() {  
    let x = 10;  
    {  
        let y = 20;  
    }  
    println!("x={}, y={}", x, y);  
}
```



# Ownership

```
// a13.rs  
fn main() {  
    let v = vec![10, 20, 30];  
    println!("{:?}", v);  
}  
// how is v deallocated?
```

# Ownership




Figure 9: Memory representation of a vector

# Ownership

```
// a14.rs  
fn fun1() {  
    let v = vec![10 ,20 ,30];  
} // how is v deallocated?  
fn main() {  
    fun1();  
}
```

# Ownership

```
// a15.rs  
fn main() {  
    let v1 = vec![10 ,20 ,30];  
    let v2 = v1;  
    println!("{:?}", v2);  
}  
// do we have a double free here?
```

# Ownership




Figure 10: Two pointers

# Ownership

```
// a15-1.rs  
fn fun(v2: Vec<i32>) {  
    println!("{:?}", v2);  
}  
fn main() {  
    let v1 = vec![10, 20, 30];  
    fun(v1);  
}  
// do we have a double free here?
```

# Ownership

```
// a16.rs  
fn main() {  
    let v1 = vec![10, 20, 30];  
    let mut v2 = v1;  
    v2.truncate(2);  
    println!("{:?}", v2);  
}  
// what happens if we try to access the  
// vector through v1?
```

# Move semantics

```
// a17.rs  
fn main() {  
    let v1 = vec![1,2,3];  
  
    let mut v2 = v1;  
    v2.truncate(2);  
    println!("{:?}", v1);  
}
```



## Move semantics

```
// a15-2.rs  
fn fun(v2: Vec<i32>) {  
    println!("{:?}", v2);  
}  
fn main() {  
    let v1 = vec![10, 20, 30];  
    fun(v1);  
    println!("{:?}", v1);  
}
```

# Move semantics

```
// a18.rs  
fn main() {  
    let a = (1, 2.3);  
    let b = a;  
    println!("{}", a);  
}
```

## Move semantics

```
// a19.rs  
fn main() {  
    let a = (1, 2.3, vec![10,20]);  
    let b = a;  
    println!("{:?}", a);  
}
```

# Memory safety without garbage collection

- ▶ Languages like Python, Java etc achieve memory safety at run time through garbage collection.
- ▶ Rust achieves memory safety at compile time by static type analysis.
- ▶ Ownership + move semantics has some interesting properties which makes them suitable for general resource management (not just memory).

# Garbage Collection

```
# a20.py  
a = [10, 20, 30]  
a.append(40)  
print a
```

# Garbage Collection

```
# a21.py  
a = [10, 20, 30]  
b = a  
b.append(40)  
print a # what does this print?
```

# Garbage Collection




Figure 11: References in Python

# Garbage Collection




Figure 12: Reference Counting



# Garbage Collection

```
# a22.py
```

```
a = [10, 20, 30]
```

```
b = a # refcount is 2
```

```
a = "hello" # refcount is 1
```

```
b = "world" # refcount drops to zero, deallocate
```

## Resource management

```
# a23.py
def read_a_line():
    f = open("/etc/passwd")
    s = f.readline()
    f.close() # close the file, release OS resources
    return s

while True:
    print read_a_line()
```

## Resource Leaks in managed languages

```
# a24.py  
def read_a_line():  
    f = open("/etc/passwd")  
    s = f.readline()  
    # No explicit "close"  
    return s  
  
while True:  
    print read_a_line()
```

## Rust means never having to close a file!

```
// a25.rs
use std::fs::File;
use std::io::Read;
fn read_whole_file() -> String {
    let mut s = String::new();
    let mut f = File::open("/etc/passwd").unwrap();
    f.read_to_string(&mut s).unwrap();
    s // return the string
}
fn main() {
    println!("{}", read_whole_file());
}
```

Read:

<http://blog.skylight.io/rust-means-never-having-to-close-a-socket/>

## Ownership / Move: Limitations

```
// a26.rs
fn vector_sum(v: Vec<i32>) -> i32 {
    //assume v is always a 3 elemnt vector
    v[0] + v[1] + v[2]
}

fn main() {
    let v = vec![1,2,3];
    let s = vector_sum(v);
    println!("{}",s);
}
```

## Ownership / Move: Limitations

```
// a27.rs  
fn vector_sum(v: Vec<i32>) -> i32 {  
    v[0] + v[1] + v[2]  
}  
fn vector_product(v: Vec<i32>) -> i32 {  
    v[0] * v[1] * v[2]  
}  
fn main() {  
    let v = vec![1,2,3];  
    let s = vector_sum(v);  
    let p = vector_product(v);  
    println!("{}",p);  
}  
// does this code compile?
```

# Immutable Borrow

```
// a28.rs  
fn vector_sum(v: &Vec<i32>) -> i32 {  
    v[0] + v[1] + v[2]  
}  
fn vector_product(v: &Vec<i32>) -> i32 {  
    v[0] * v[1] * v[2]  
}  
fn main() {  
    let v = vec![1,2,3];  
    let s = vector_sum(&v);  
    let p = vector_product(&v);  
    println!("v={:?}, s={}, p={}", v, s, p);  
}
```

# Immutable Borrow

```
// a29.rs  
fn main() {  
    let v = vec![1,2,3];  
    let t1 = &v;  
    let t2 = &v;  
    println!("{}", t1[0], t2[0], v[0]);  
}  
// any number of immutable borrows are ok!
```



# Immutable Borrow

```
// a30.rs  
fn change(t1: &Vec<i32>) {  
    t1[0] = 10;  
}  
fn main() {  
    let mut v = vec![1,2,3];  
    change(&v);  
}  
// Does the program compile?
```

# Mutable Borrow

```
// a31.rs
fn change(t1: &mut Vec<i32>) {
    t1[0] = 10;
}

fn main() {
    let mut v = vec![1,2,3];
    change(&mut v);
    println!("{:?}", v);
}
```

## A use-after-free bug

```
// a32.c
#include <stdlib.h>
int main()
{
    char *p = malloc(10 * sizeof(char));
    char *q;

    q = p + 2;
    free(p);
    *q = 'A'; // bug!
    return 0;
}
```

## Vector allocation in Rust

```
// a33.rs  
fn main() {  
    let mut a = vec![];  
    a.push(1); a.push(2);  
    a.push(3); a.push(4);  
    a.push(5);  
  
    println!("{:?}", a);  
}
```

# Vector allocation in Rust/C++




Figure 13: Growing a vector

## A use-after-free bug in C++

```
// a33-1.cpp
#include <vector>
#include <iostream>
using namespace std;
int main()
{
    vector<int> v;
    int *p;
    v.push_back(1);
    p = &v[0];
    v.push_back(2);
    *p = 100; // bug!
    cout << v[0] << endl;
}
```

## A use-after-free bug in Rust?

```
// a34.rs  
fn main() {  
    let mut v = vec![10, 20, 30, 40];  
    let p1 = &v[1];  
    v.push(50);  
    // bug if we try to use p1  
    // does this code compile?  
}
```

# Borrowing Rules

- ▶ Any number of immutable borrows can co-exist.
- ▶ A mutable borrow can not co-exist with other mutable or immutable borrows.
- ▶ The “borrow checker” checks violations of these rules at compile time.



## Borrow checker limitations

- ▶ The borrow checker gives you safety by rejecting ALL unsafe programs.
- ▶ But it is not perfect in the sense it rejects safe programs also; “fighting the borrow checker” is a common sporting activity among Rust programmers :)
- ▶ There are plans to improve the situation:  
<http://smallcultfollowing.com/babysteps/blog/2017/03/01/nested-method-calls-via-two-phase-borrowing/>

## Borrow checker limitations - an example

```
// a35.rs  
fn main() {  
    let mut v = vec![10,20,30];  
    v.push(v.len());  
}  
// this will not compile
```

## Borrow checker limitations - an example

```
// a36.rs  
// Same as a35.rs  
fn main() {  
    let mut v = vec![10,20,30];  
    let tmp0 = &v;  
    let tmp1 = &mut v;  
    let tmp2 = Vec::len(tmp0); //v.len()  
  
    Vec::push(tmp1, tmp2); // v.push(tmp2)  
}
```

# Lifetimes

```
// a37.rs  
fn main() {  
    let ref1: &Vec<i32>;  
    {  
        let v = vec![1, 2, 3];  
        ref1 = &v;  
    }  
    // v gets deallocated as it goes out of  
    // the scope. What about ref1? Do we have  
    // a "dangling pointer" here?  
}
```

# Lifetimes

```
// a38.rs
fn foo() -> Vec<i32> {
    let v = vec![1, 2, 3];
    v // transfer ownership to caller
}

fn main() {
    let p = foo();
    println!("{:?}", p);
}
```

# Lifetimes

```
// a39.rs
fn foo() -> &Vec<i32> {
    let v = vec![1, 2, 3];
    &v // Will this compile?
}

fn main() {
    let p = foo();
}
```

## Explicit Lifetime Annotations

```
// a40.rs
fn foo(v1: &Vec<i32>, v2: &Vec<i32>) -> &i32 {
    &v1[0]
}

fn main() {
    let v1 = vec![1, 2, 3];
    let p:&i32;
    {
        let v2 = vec![4, 5, 6];
        p = foo(&v1, &v2);
        // How does the compiler know, just by looking at
        // the signature of "foo", that the reference
        // returned by "foo" will live as long as "p"?
    }
}
```

## Explicit Lifetime Annotations

```
// a41.rs
fn foo<'a, 'b>(v1: &'a Vec<i32>,
               v2: &'b Vec<i32>) -> &'a i32 {

    &v1[0]
}

fn main() {
    let v1 = vec![1, 2, 3];
    let p:&i32;
    {
        let v2 = vec![4, 5, 6];
        p = foo(&v1, &v2);
    }
}
```



## Explicit Lifetime Annotations

```
// a42.rs
fn foo<'a, 'b>(v1: &'a Vec<i32>,
               v2: &'b Vec<i32>) -> &'b i32 {

    &v2[0]
}

fn main() {
    let v1 = vec![1, 2, 3];
    let p:&i32;
    {
        let v2 = vec![4, 5, 6];
        p = foo(&v1, &v2);
    }
}
```

# Unsafe

```
// a43.rs
fn main() {
    // a is a "raw" pointer initialized to 0
    let a: *mut u32 = 0 as *mut u32;

    *a = 0;
}
```

# Unsafe

```
// a44.rs  
fn main() {  
    let a: *mut u32 = 0 as *mut u32;  
  
    unsafe {  
        *a = 0;  
    }  
}
```

## End of Part 1

- ▶ What we have seen so far is the “core” of Rust, these are the ideas which make Rust unique!
- ▶ Most of the other “interesting” ideas are borrowed from statically typed functional programming languages (like ML). (The first Rust compiler was written in Ocaml).

# End of Part 1



**llogiq**  
@llogiq



Follow



#rustlang is a very strange place  
sans null deref nor data race  
it has its own styles  
but once it compiles  
it will not blow up in your face

RETWEETS

3

LIKES

7



3:41 am - 2 Mar 2017



1



3



7

Figure 14: rustpoem

# Zero Cost Abstractions

```
// a45.rs
const N: u64 = 1000000000;
fn main() {
    let r = (0..N)
        .map(|x| x + 1)
        .fold(0, |sum, i| sum+i);

    println!("{}", r);
}
// Compile with optimizations enabled:
// rustc -O a45.rs
```

## Zero cost abstractions

Here is part of the assembly language code produced by the compiler for a45.rs:

```
.Ltmp0:
    .cfi_def_cfa_offset 80
    movabsq $50000000005000000000, %rax
    movq    %rax, (%rsp)
    leaq    (%rsp), %rax
    movq    %rax, 8(%rsp)
```

Looks like the expression has been evaluated fully at compile time itself!

Here is the commandline used to produce the above output:

```
rustc -O a45.rs --emit=asm
```

# Zero cost abstractions

- ▶ You can write confidently using all the high-level abstractions the language has to offer.
- ▶ Your code will almost always be as fast as hand-coded low level C!



## Sum Types and Pattern Matching

```
// a46.rs
enum Color {
    Red,
    Green,
    Blue,
}

use Color::*;

fn main() {
    let c = Red;
    match c {
        Red => println!("color is Red!"),
        Green => println!("color is Green!"),
        Blue => println!("color is Blue!")
    }
}
```

## Sum Types and Pattern Matching

```
// a47.rs
#[derive(Debug)]
enum Shape {
    Circle(u32),
    Square (u32),
    Rectangle {ht: u32, wid: u32},
}

use Shape::*;
fn main() {
    let s1 = Circle(10);
    let s2 = Square(5);
    let s3 = Rectangle {ht: 10, wid: 2};

    println!("{:?}", s3);
}
```

## Pattern matching is exhaustive

```
// a48.rs
#[derive(Debug)]
enum Shape {
    Circle(f64),
    Square (f64),
    Rectangle {ht: f64, wid: f64},
}

use Shape::*;
fn area(s: Shape) -> f64 {
    match s {
        Circle(x) => 3.14 * x * x,
        Rectangle {ht: x, wid: y} => x * y,
    } // bug!
}

fn main() {
    let s1 = Circle(10.0);
    println!("{}", area(s1));
}
```

# The Option type

```
// a49.rs  
fn main() {  
    let mut a = vec![10];  
    let b = a.pop();  
    let c = a.pop();  
  
    let d = b + 1; // does it compile?  
}
```

# The Option type

```
// a50.rs  
fn main() {  
    let mut a = vec![10];  
    let b = a.pop();  
    let c = a.pop();  
  
    println!("b = {:?}, c = {:?}", b, c);  
}
```

## The Option type

```
// a51.rs  
fn main() {  
    let mut a = vec![10];  
    let b = a.pop();  
  
    match b {  
        Some(x) => println!("pop: {}", x),  
        None => println!("empty stack"),  
    }  
}
```

## The Option type

```
// a52.rs  
fn main() {  
    let mut a = vec![10];  
    let b = a.pop();  
  
    println!("{}", b.unwrap());  
    let c = a.pop();  
    println!("{}", c.unwrap());  
}
```