

The Rust Programming Language

Pramode C.E

March 2, 2017

A language that doesn't affect the way you think about programming, is not worth knowing.

Alan Perlis.

Why is Rust so exciting?

- ▶ Safe low-level systems programming
- ▶ Memory safety without garbage collection
- ▶ High-level abstractions (influenced by statically typed functional programming languages like ML) without run-time overhead
- ▶ Concurrency without data races

Why not just use C/C++?

The second big thing each student should learn is: How can I avoid being burned by C's numerous and severe shortcomings? This is a development environment where only the paranoid can survive.

<http://blog.regehr.org/archives/1393>

A bit of Rust history

- ▶ Started by Graydon Hoare as a personal project in 2006
- ▶ Mozilla foundation started sponsoring Rust in 2010
- ▶ Rust 1.0 released in May, 2015
- ▶ Regular six week release cycles

Core language features

- ▶ Memory safety without garbage collection
 - ▶ Ownership
 - ▶ Move Semantics
 - ▶ Borrowing and lifetimes
- ▶ Type Inference
- ▶ Algebraic Data Types (Sum and Product types)
- ▶ Exhaustive Pattern Matching
- ▶ Trait-based generics
- ▶ Iterators
- ▶ Zero Cost Abstractions
- ▶ Concurrency without data races
- ▶ Efficient C bindings, minimal runtime

Structure of this workshop

- ▶ Understanding the problems with C/C++
- ▶ Understanding Ownership, Borrow, Move semantics and Lifetimes
- ▶ Other features (depending on availability of time)

The Stack

```
// a0.c
```

```
int main()
{
    int a = 1, b = 2;
    return 0;
}
```


The Stack




Figure 1: The C stack

Buffer overflow

```
// a1.c
int main()
{
    int i=1;
    int a[4];
    int j=2;

    a[4] = 5; // bug
    a[5] = 6; // bug
    a[10000] = 7; // bug
}
```

Buffer overflow




Figure 2: Buffer overflow diagram

Pointers in C

```
// a2.c  
int main()  
{  
    int a = 10;  
    int *b;  
  
    b = &a;  
    *b = 20;  
}
```

Pointers in C




Figure 3: How pointer variables are stored in memory

Stack Frames - deallocations and allocations

```
// a3.c  
void fun2() { int e=5, f=6; }  
void fun1() { int c=3, d=4; }  
  
int main()  
{  
    int a=1, b=2;  
    fun1(); fun2();  
    return 0;  
}
```

Stack Frames - allocations and deallocations




Figure 4: Multiple stack frames

Stack Frames - allocations and deallocations




Figure 5: Multiple stack frames

Dangling Pointers

```
// a4.c
void fun2() { int m = 1; int n = 2; }
int* fun1() {
    int *p; int q = 0;
    p = &q; return p; // bug
}

int main() {
    int *a, b; a = fun1();
    *a = 10; fun2();
    b = *a;
}
```

Dangling pointers




Figure 6: Dangling pointer

Null pointer dereferencing

```
// a6.c
#include <strings.h>
#include <stdio.h>
int main()
{
    char s[100] = "hello";
    char *p;
    p = index(s, 'f');
    *p = 'a'; // bug!
    return 0;
}
```

Heap allocation - malloc and free

```
// a7.c
#include <stdlib.h>
void fun()
{
    char *c;
    c = malloc(10*sizeof(char));
    /* do some stuff here */
    free(c);
}
int main()
{
    fun();
}
```

Heap allocation - malloc and free




Figure 7: A stack location pointing to a heap location

Memory leaks

```
// a8.c
#include <stdlib.h>
void fun()
{
    char *c;
    c = malloc(10*sizeof(char));
    /* do some stuff here */
}
int main()
{
    fun(); // bug! memory leak.
}
```

Use-after-free

```
// a9.c
#include <stdlib.h>
void fun(char *t) {
    /* do some stuff here */
    free(t);
}
int main() {
    char *c;
    c = malloc(10 * sizeof(char));
    fun(c);
    c[0] = 'A'; //bug! user-after-free
}
```

Double free

```
// a10.c
#include <stdlib.h>
void fun(char *t) {
    /* do some stuff here */
    free(t);
}
int main() {
    char *c;
    c = malloc(10 * sizeof(char));
    fun(c);
    free(c); //bug! double free
}
```


Undefined behaviours and optimization

```
// a11.c
#include <limits.h>
#include <stdio.h>
// compile the code with optimization (-O3) and without
int main() {
    int c = INT_MAX;
    if (c+1 < c)
        printf("hello\n");
    printf("%d\n", c+1);
}
```

Undefined behaviours




Figure 8: https://www.explainxkcd.com/wiki/images/c/cc/cant_sleep.png

Undefined behaviours

When tools like the bounds checking GCC, Purify, Valgrind, etc. first showed up, it was interesting to run a random UNIX utility under them. The output of the checker showed that these utility programs, despite working perfectly well, executed a ton of memory safety errors such as use of uninitialized data, accesses beyond the ends of arrays, etc. Just running grep or whatever would cause tens or hundreds of these errors to happen.

From: <http://blog.regehr.org/archives/226>

Undefined behaviours

More and more, I'm starting to wonder how safety-critical code can continue being written in C.

A comment on: <http://blog.regehr.org/archives/232>

Hello, world!

```
// a12.rs  
fn main() {  
    println!("hello, world!");  
}
```

Compile: `rustc a12.rs`

Run: `./a12`

Type inference

```
// a12-1.rs
fn sqr(x: i32) -> i32 {
    let y = x * x; // type inferred
    y
}
fn main() {
    let t1 = sqr(10); // type inferred
    let t2:i32 = sqr(20);
    println!("sqr 10 = {}, sqr 20 ={}", t1, t2);
}
```

Immutability

```
// a12-2.rs
fn main() {
    let x = 0; // x is immutable
    let mut y = 1;
    x = x + 1; // does not work
    y = y + 1;
}
```

Scope

```
// a12-3.rs  
fn main() {  
    let x = 10;  
    {  
        let y = 20;  
    }  
    println!("x={}, y={}", x, y);  
}
```


Ownership

```
// a13.rs  
fn main() {  
    let v = vec![10, 20, 30];  
    println!("{:?}", v);  
}  
// how is v deallocated?
```

Ownership




Figure 9: Memory representation of a vector

Ownership

```
// a14.rs  
fn fun1() {  
    let v = vec![10 ,20 ,30];  
} // how is v deallocated?  
fn main() {  
    fun1();  
}
```

Ownership

```
// a15.rs  
fn main() {  
    let v1 = vec![10 ,20 ,30];  
    let v2 = v1;  
    println!("{:?}", v2);  
}  
// do we have a double free here?
```

Ownership




Figure 10: Two pointers

Ownership

```
// a15-1.rs  
fn fun(v2: Vec<i32>) {  
    println!("{:?}", v2);  
}  
fn main() {  
    let v1 = vec![10, 20, 30];  
    fun(v1);  
}  
// do we have a double free here?
```

Ownership

```
// a16.rs  
fn main() {  
    let v1 = vec![10, 20, 30];  
    let mut v2 = v1;  
    v2.truncate(2);  
    println!("{:?}", v2);  
}  
// what happens if we try to acces the  
// vector through v1?
```

Move semantics

```
// a17.rs  
fn main() {  
    let v1 = vec![1,2,3];  
  
    let mut v2 = v1;  
    v2.truncate(2);  
    println!("{:?}", v1);  
}
```


Move semantics

```
// a15-2.rs  
fn fun(v2: Vec<i32>) {  
    println!("{:?}", v2);  
}  
fn main() {  
    let v1 = vec![10, 20, 30];  
    fun(v1);  
    println!("{:?}", v1);  
}
```

Move semantics

```
// a18.rs  
fn main() {  
    let a = (1, 2.3);  
    let b = a;  
    println!("{}", a);  
}
```

Move semantics

```
// a19.rs  
fn main() {  
    let a = (1, 2.3, vec![10,20]);  
    let b = a;  
    println!("{:?}", a);  
}
```

Memory safety without garbage collection

- ▶ Languages like Python, Java etc achieve memory safety at run time through garbage collection.
- ▶ Rust achieves memory safety at compile time by static type analysis.
- ▶ Ownership + move semantics has some interesting properties which makes them suitable for general resource management (not just memory).

Garbage Collection

```
# a20.py  
a = [10, 20, 30]  
a.append(40)  
print a
```

Garbage Collection

```
# a21.py  
a = [10, 20, 30]  
b = a  
b.append(40)  
print a # what does this print?
```

Garbage Collection




Figure 11: References in Python

Garbage Collection




Figure 12: Reference Counting

Garbage Collection

```
# a22.py
```

```
a = [10, 20, 30]
```

```
b = a # refcount is 2
```

```
a = "hello" # refcount is 1
```

```
b = "world" # refcount drops to zero, deallocate
```

Resource management

```
# a23.py
def read_a_line():
    f = open("/etc/passwd")
    s = f.readline()
    f.close() # close the file, release OS resources
    return s

while True:
    print read_a_line()
```

Resource Leaks in managed languages

```
# a24.py  
def read_a_line():  
    f = open("/etc/passwd")  
    s = f.readline()  
    # No explicit "close"  
    return s  
  
while True:  
    print read_a_line()
```

Rust means never having to close a file!

```
// a25.rs
use std::fs::File;
use std::io::Read;
fn read_whole_file() -> String {
    let mut s = String::new();
    let mut f = File::open("/etc/passwd").unwrap();
    f.read_to_string(&mut s).unwrap();
    s // return the string
}
fn main() {
    println!("{}", read_whole_file());
}
```

Read:

<http://blog.skylight.io/rust-means-never-having-to-close-a-socket/>

Ownership / Move: Limitations

```
// a26.rs
fn vector_sum(v: Vec<i32>) -> i32 {
    //assume v is always a 3 elemnt vector
    v[0] + v[1] + v[2]
}

fn main() {
    let v = vec![1,2,3];
    let s = vector_sum(v);
    println!("{}",s);
}
```

Ownership / Move: Limitations

```
// a27.rs  
fn vector_sum(v: Vec<i32>) -> i32 {  
    v[0] + v[1] + v[2]  
}  
fn vector_product(v: Vec<i32>) -> i32 {  
    v[0] * v[1] * v[2]  
}  
fn main() {  
    let v = vec![1,2,3];  
    let s = vector_sum(v);  
    let p = vector_product(v);  
    println!("{}",p);  
}  
// does this code compile?
```

Immutable Borrow

```
// a28.rs  
fn vector_sum(v: &Vec<i32>) -> i32 {  
    v[0] + v[1] + v[2]  
}  
fn vector_product(v: &Vec<i32>) -> i32 {  
    v[0] * v[1] * v[2]  
}  
fn main() {  
    let v = vec![1,2,3];  
    let s = vector_sum(&v);  
    let p = vector_product(&v);  
    println!("v={:?}, s={}, p={}", v, s, p);  
}
```

Immutable Borrow

```
// a29.rs  
fn main() {  
    let v = vec![1,2,3];  
    let t1 = &v;  
    let t2 = &v;  
    println!("{}", t1[0], t2[0], v[0]);  
}  
// any number of immutable borrows are ok!
```


Immutable Borrow

```
// a30.rs  
fn change(t1: &Vec<i32>) {  
    t1[0] = 10;  
}  
fn main() {  
    let mut v = vec![1,2,3];  
    change(&v);  
}  
// Does the program compile?
```

Mutable Borrow

```
// a31.rs  
fn change(t1: &mut Vec<i32>) {  
    t1[0] = 10;  
}  
fn main() {  
    let mut v = vec![1,2,3];  
    change(&mut v);  
    println!("{:?}", v);  
}
```

A use-after-free bug

```
// a32.c
#include <stdlib.h>
int main()
{
    char *p = malloc(10 * sizeof(char));
    char *q;

    q = p + 2;
    free(p);
    *q = 'A'; // bug!
    return 0;
}
```

Vector allocation in Rust

```
// a33.rs  
fn main() {  
    let mut a = vec![];  
    a.push(1); a.push(2);  
    a.push(3); a.push(4);  
    a.push(5);  
  
    println!("{:?}", a);  
}
```

Vector allocation in Rust




Figure 13: Growing a vector

A use-after-free bug in Rust?

```
fn main() {  
    let mut v = vec![10, 20, 30, 40];  
    let p1 = &v[1];  
    v.push(50);  
    // bug if we try to use p1  
    // does this code compile?  
}
```

Borrowing Rules

- ▶ Any number of immutable borrows can co-exist.
- ▶ A mutable borrow can not co-exist with other mutable or immutable borrows.
- ▶ The “borrow checker” checks violations of these rules at compile time.

Conclusion



llogiq
@llogiq

 Follow



#rustlang is a very strange place
sans null deref nor data race
it has its own styles
but once it compiles
it will not blow up in your face

RETWEETS

3

LIKES

7



3:41 am - 2 Mar 2017



1



3



7

Figure 14: rustpoem